



Nix for Polyglot, Reproducible Data Science Workflows

Bruno Rodrigues 

Ministry of Research and Higher education, Luxembourg

Philipp Baumann 

Alliance SwissPass, Switzerland

Abstract

Reproducible analysis requires more than clean, well-documented code; it demands the precise management of software dependencies, computational environments, and workflow execution. While researchers can combine tools like **Docker**, **renv** for R packages (or **uv** for Python packages) and **Make** to address these needs, the resulting toolchain is often complex and fragile, requiring the coordination of multiple disparate systems. This paper introduces a unified framework built on the **Nix** package manager, made accessible through two R packages. The first, **rix**, generates declarative **Nix** expressions to define version-pinned, reproducible environments that encompass R, Python, Julia, and system-level dependencies. The second, **rixxpress** (and its Python port, **ryxxpress**), leverages these environments to orchestrate polyglot pipelines where each computational step runs in a hermetically sealed, language-specific environment with automatic caching. By treating the entire computational environment as code, this integrated approach ensures a high degree of reproducibility, supports collaboration across heterogeneous systems, and helps guarantee that analyses remain executable long into the future.

Keywords: reproducibility, R, Python, Julia, Nix.

1. Introduction: Reproducibility is also about software

Peng (2011) introduced the concept of reproducibility as a *continuum*. At one end lies the least reproducible state, where only a paper describing the study is available. Reproducibility improves when authors share the original source code, improves further when they include the underlying data, and reaches its highest level when what Roger Peng called *linked and executable code and data* are provided.

By *linked and executable code and data*, Peng referred to compiled source code and runnable scripts. In this paper, we interpret this notion more broadly as the *computational environment*:

the complete set of software required to execute an analysis. Here too, a continuum exists. At the minimal end, authors might only name the main software used—say, the R programming language (R Core Team 2021). More careful authors might also specify the version of R, or list the additional packages and their versions. Rarely, however, do authors specify the operating system on which the analysis was performed, even though differences in operating systems can lead to divergent results when using the same code and software versions, as shown by Bhandari Neupane, Neupane, Luo, Yoshida, Sun, and Williams (2019). It is even less common for authors to provide step-by-step installation instructions for the required software stack, an omission often driven by institutional constraints. Journals, for instance, can inadvertently hinder reproducibility by imposing strict page or word limits that leave no room for the necessary technical documentation, discouraging thoroughness in the name of brevity.

Even when such instructions are given, they often fail across different platforms or versions of the same platform. This lack of portability not only hinders reproducibility but also complicates everyday research workflows. Researchers using multiple machines must recreate environments consistently, and collaborators must share identical computational setups to avoid inconsistencies.

Finally, once the execution environment is correctly configured, additional clarity is needed on how to *run* the project itself. Which packages should be loaded first? Which scripts should be executed, and in what order? Without clear documentation or automated orchestration, these operational details become yet another barrier to reproducibility.

This paper focuses on two critical but often overlooked aspects of reproducibility: *computational environment management* and *workflow orchestration*. We present a comprehensive framework addressing both challenges using the Nix package manager (Dolstra, De Jonge, and Visser 2004), making it accessible to researchers through two R packages: **rix** and **rixpress**. Before introducing these packages, we survey existing tools and their limitations to contextualise our contributions.

A range of tools now exist to help researchers approach the gold standard of full reproducibility, or to consistently deploy the same development environment across multiple machines. Let us first consider the most basic step in this process: listing the software used. In R, the `sessionInfo()` function provides a concise summary of the software environment, including the R version, platform details, and all loaded packages. Its output can be saved to a file and included as part of a study’s reproducibility record. Below is an example output from `sessionInfo()`:

```
R> sessionInfo()

R version 4.3.2 (2023-10-31)
Platform: aarch64-unknown-linux-gnu (64-bit)
Running under: Ubuntu 22.04.3 LTS

Matrix products: default
BLAS:   /usr/lib/aarch64-linux-gnu/openblas-pthread/libblas.so.3
LAPACK: /usr/lib/aarch64-linux-gnu/openblas-pthread/libopenblas[...]
```

locale:

```
LC_CTYPE=en_US.UTF-8      LC_NUMERIC=C
LC_TIME=en_US.UTF-8      LC_COLLATE=en_US.UTF-8
...
```

attached base packages:

```
stats      graphics  grDevices  utils      datasets  methods   base
```

other attached packages:

```
nnet_7.3-19  mgcv_1.9-0  nlme_3.1-163
```

loaded via a namespace (and not attached):

```
compiler_4.3.2  Matrix_1.6-1.1  tools_4.3.2      splines_4.3.2
grid_4.3.2      lattice_0.21-9
```

When an author includes this information, others attempting to reproduce the study (future readers, collaborators, or even the author at a later time) can see which version of R and which packages (with their versions) were used. However, reproducing the environment still requires manually installing the correct package versions — a process that becomes difficult when packages depend on system-level libraries. For example, the **sf** package requires **GDAL**, **GEOS**, and **PROJ** libraries. Installing and configuring these dependencies varies substantially across operating systems and can be prohibitively difficult on some platforms.

A more robust approach than simply listing package versions is to use tools that define per-project environments. This is common for Python users, who have many tools at their disposal, such as the built-in *virtual environments* module, or third-party tools like **uv**, which manages both project-specific interpreters and package libraries using a lock file mechanism.

In the R ecosystem, virtual environments are not built-in as they are in Python, however tools like **renv** provide equivalent functionality. **renv** captures the project’s software state and writes it to a lock file (**renv.lock**), which includes the exact versions of R and all required packages. This lockfile serves as a blueprint for restoring the environment automatically, ensuring others can recreate the same setup with minimal effort.

However, both **renv** and **uv** have limitations. While **uv** manages Python interpreter versions directly, **renv** does not restore the version of R itself—installing the correct version must be done separately using tools like **rig** (R Infrastructure 2023). More critically, neither handles system-level dependencies. If **sf** requires **GDAL** version 3.0 but the system has version 2.4 installed, neither tool can resolve this conflict. Users must manually install system libraries, and the process differs across operating systems.

Pre-compiled binaries simplify installation but do not eliminate the need for system libraries required at runtime. Other R packages such as **groundhog** (Simonsohn and Gruson 2023) and **rang** (hong Chan and Schoch 2023) enable date-based package installation, while the Posit Package Manager provides dated CRAN snapshots. However, like **renv**, these focus on R packages and do not address system-level dependencies or R version management.

These tools represent significant progress in reproducibility within their ecosystems, but they share a common limitation: none handles system-level dependencies or polyglot environments. A project using R, Python (Van Rossum and Drake 2009), and system tools like **Quarto** or

L^AT_EX requires coordinating multiple package managers (**renv** for R, **uv** for Python, manual installation for system tools), each with its own configuration and potential conflicts.

The most comprehensive approach to date is containerisation with **Docker**. **Docker** packages a *data product* together with all its dependencies—including the operating system, system libraries, programming languages, and packages—into a self-contained image. Once built, the analysis can be executed inside a *container*, a running instance of the image.

The strength of **Docker** lies in its ability to bundle not only the correct versions of R and its packages but also system-level dependencies. Since **Docker** images are minimal Linux systems, they include the libraries required by packages such as **sf**, ensuring that future replicators have access to the same environment. Moreover, these images can be shared via registries like Docker Hub, enabling reproducibility across systems.

The Rocker project (Boettiger and Eddelbuettel 2017) provides pre-built **Docker** images for the R community, offering specific R versions and, in some cases, preinstalled packages. These serve as convenient bases for reproducible analysis environments.

Despite its strengths, **Docker** presents several challenges:

1. Poor interactive support: While containers can be modified at runtime, changes are lost unless explicitly saved. Running graphical applications like RStudio Desktop requires complex X11 configuration that works reliably only on Linux. Web-based IDEs such as RStudio Server are alternatives but have their own limitations.
2. Steep learning curve: Writing reliable **Dockerfiles** requires familiarity with Linux system administration. To ensure reproducibility, **Dockerfiles** should reference image *digests* rather than *tags*, but this is rarely done. Packages like **dockerfiler** (Fay, Guyader, Parry, and Rochette 2024) help automate this but do not remove the need for basic Linux knowledge.
3. Post-hoc reproducibility: Analyses are often developed interactively and containerised only afterward, achieving post-hoc reproducibility rather than consistent environments during development.
4. Single-language focus: While **Docker** can support polyglot setups, orchestrating them requires manual coordination. There is no standard way to specify that one step runs in Python and another in R.
5. Incomplete reproducibility: Reproducibility is best viewed as a spectrum rather than a binary property (Dellaiera 2024). **Docker** achieves strong run-time reproducibility—consistency across systems (space)—but weak build-time reproducibility, or consistency over time. Mutable base images and non-deterministic commands such as **apt-get update** introduce temporal drift: rebuilding the same **Dockerfile** at different times can yield different results. Even with pinned base images, container checksums can vary across builds and machines. Thus, while **Docker** provides strong run-time reproducibility, achieving deterministic artefacts requires pinning image digests, freezing dependencies, and using deterministic build systems such as **Nix** for full control over both time and space (Malka, Zacchiroli, and Zimmermann 2024).

Despite these shortcomings, **Docker** remains valuable for specific use cases: production deployment, CI/CD integration, and cross-platform distribution. The two approaches are not

mutually exclusive. **Nix** can be installed and used inside **Docker**, combining **Nix**'s deterministic builds with **Docker**'s deployment infrastructure. In this workflow, **Nix** handles environment reproducibility during the build of the image, while **Docker** serves as the distribution mechanism. This addresses **Docker**'s build-time reproducibility issues while maintaining its deployment advantages. However, for research workflows where interactive development is primary and **Docker**'s complexity is a barrier, **Nix** alone may suffice.

Thankfully, it is entirely possible to combine both **Docker** and **Nix**: **Nix** to setup reproducible environments and **Docker** for deployment across the many different already existing infrastructures.

Even with a perfectly reproducible environment, researchers face another challenge: reliably executing the analysis workflow. Which scripts run first? What are the dependencies between steps? Manual execution is error-prone and poorly documented. Build automation tools like **Make** address this by defining analyses as ordered, reproducible steps.

Within R, the **targets** package (Landau 2021) provides a modern, declarative approach to workflow management. It tracks dependencies, caches results, and only recomputes affected steps, improving efficiency for complex analyses. However, **targets** operates within a single R session. For truly polyglot pipelines—where steps may require incompatible dependencies or different language versions—this becomes problematic. Running **targets** inside a **Docker** container ensures reproducibility but forces the entire pipeline to share one environment.

McDermott (2021) provides an excellent example of best practice. The accompanying repository¹ demonstrates:

- Package versions recorded in an `renv.lock` file;
- A `Makefile` automating the full analysis;
- A `Dockerfile` providing the computational environment.

Achieving this required mastering multiple complex tools: **renv** for package management, **Docker** for containerisation, and **Make** for orchestration. The complexity increases for polyglot projects, which must manage multiple ecosystems simultaneously.

Researchers therefore need a tool that:

1. Manages complete environments including languages, packages, and system dependencies;
2. Supports multiple languages natively;
3. Works interactively without container complexity;
4. Provides step-level isolation for different environments;
5. Ensures bit-for-bit reproducibility through deterministic builds;
6. Remains simple enough for researchers without systems administration expertise.

The **Nix** package manager provides these capabilities. It ensures reproducible installation by deploying *component closures*—packages bundled with all their direct and transitive dependencies (Dolstra *et al.* 2004). Think of installing a package as packing for a trip: traditional package managers assume you'll find essentials at your destination, while **Nix** packs everything you need. This produces self-contained software environments that behave identically across machines and over time.

¹See <https://github.com/grantmcdermott/skeptic-priors>

Nix can replace **Docker** for isolation, **renv** for package management, and even **Make** for orchestration—all within a single, unified framework. However, it has a steep learning curve, with its own functional language for declaratively defining software builds and configurations. While this ensures reproducibility, it also creates a barrier to adoption.

To make **Nix** accessible to researchers, we developed two R packages:

- **rix** generates **Nix** expressions from intuitive R function calls, eliminating the need to learn the Nix language. It handles environment definitions including R, Python, Julia (Bezanson, Edelman, Karpinski, and Shah 2017), system tools, and dependencies.
- **rixpress** orchestrates polyglot analytical pipelines using **Nix** as the build engine. Each step runs in its own hermetic environment with automatic caching and dependency tracking, enabling true polyglot workflows across R, Python, and Julia. Python users can use **ryxpress**, a direct port.

Together, these packages provide a single framework for managing reproducible, polyglot environments and executing complex workflows with step-level isolation. While **Nix** remains complex for advanced use, **rix** and **rixpress** abstract this complexity, making deep reproducibility accessible without systems administration expertise.

The remainder of this paper proceeds as follows. Section 2 introduces **Nix** and explains how its functional model enables reproducibility. Section 3 presents **rix** and demonstrates environment definition. Section 4 discusses the **rstats-on-nix** fork of the package repository. Section 5 introduces **rixpress** and demonstrates polyglot pipeline orchestration. Section 6 concludes with discussion of future directions.

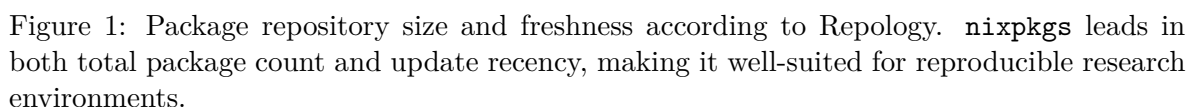
2. The Nix package manager

Nix is a powerful, cross-platform package manager designed for installing and building software in a fully reproducible and reliable manner. Unlike traditional package managers, **Nix** emphasises immutable infrastructure and a functional approach, which significantly enhances the consistency of computational environments. As of writing, the primary **Nix** package collection, **nixpkgs**, provides access to over 120,000 packages, including nearly all of CRAN and Bioconductor. According to Repology, a service that tracks package repositories across distributions, **nixpkgs** is both the largest and among the most up-to-date package repositories available (see Figure Figure 1).

This extensive coverage allows users to install not only R itself but also all its required packages and system-level dependencies for any given project.

While **Nix** is the default package manager for the NixOS Linux distribution, it can be installed as a standalone tool on other Linux distributions, macOS, and effectively on Windows via the Windows Subsystem for Linux (WSL2), treating these as equivalent platforms in practice.

A key advantage of using **Nix** to install R packages, as opposed to the standard `install.packages()` function, lies in its comprehensive dependency management. **Nix** ensures that all dependencies of each package are installed, irrespective of whether they are other R packages or underlying system libraries.



This concept is central to **Nix**’s design, where packages are referred to as *component closures*. As [Dolstra *et al.* \(2004\)](#) explains:

The core of **Nix**’s reproducibility lies in its functional package management paradigm and its unique approach to builds. When installing software with **Nix**, it evaluates an *expression* written in the Nix language. These expressions define *derivations*, which are declarative

blueprints describing how to build a package. A derivation specifies all inputs: source code, build commands, and crucially, all direct and transitive dependencies. This declarative nature ensures that every build process is fully specified and independent of the environment where it's executed. Furthermore, all **Nix** expressions for official packages are hosted in the `nixpkgs` GitHub repository. By referencing a specific commit of `nixpkgs` (a process known as *pinning a revision*) users guarantee that all packages installed are derived from the exact same set of build instructions. This means that if an environment is built today using a pinned revision, it will produce the identical set of software versions and configurations if rebuilt tomorrow, next year, or on a different machine, as long as the `nixpkgs` repository remains accessible.

Pinning is essential, but it is not the only mechanism through which **Nix** achieves reproducibility. **Nix** adopts principles from functional programming, conceptualising software builds as *pure functions*. In mathematics, a pure function is defined as one that always produces the same output for a given input—for instance, if $f(x) = x^2$, then $f(2) = 4$ invariably holds. Analogously, **Nix** ensures that, given identical inputs (such as source code, dependencies, and build instructions), the resulting output (the built package) remains identical, irrespective of the environment or time of execution. In this sense, **Nix** functions as a *functional package manager*, applying the deterministic properties of pure functions to software construction.

This is achieved by disallowing side effects and ensuring that builds operate within hermetic (isolated) environments, with no hidden dependencies on global system state. The principles of functional package management (hermetic builds, immutable artifacts, and declarative dependencies) are designed to produce deterministic outputs. The effectiveness of this model was recently validated in the most extensive study of build reproducibility to date. The effectiveness of this model was recently validated at unprecedented scale by [Malka, Zacchiroli, and Zimmermann \(2025\)](#), who rebuilt 709,816 packages from historical `nixpkgs` snapshots (2017–2023). They found rebuildability rates exceeding 99% and bitwise reproducibility rates between 69% and 91%, with an upward trend over time. Importantly, they identified that about 15% of unreproducible builds failed due to embedded build timestamps, which a solvable issue rather than a fundamental limitation.

While this functional approach greatly enhances reproducibility, it can sometimes introduce complexity for package maintainers, especially for software that needs to download external assets during installation (e.g., certain Bioconductor packages). For end-users, however, this complexity is largely abstracted away thanks to both **rix** and **rixpress**.

Additionally, **Nix** supports multiple versions or *variants* of a package on the same system. **Nix** installs all packages into a centralised, content-addressed directory called the **Nix** store (typically `/nix/store`). Each package gets a unique subdirectory whose name includes a hash of all its inputs: source code, dependencies, build instructions. This means `/nix/store/abc123-dplyr-1.0.0` and `/nix/store/xyz789-dplyr-1.1.0` can coexist without conflicts.

This isolation ensures that projects always use their intended software versions, allowing multiple versions of R to coexist and preventing “dependency hell” and global environment pollution.

For a more in-depth technical discussion of **Nix**'s design principles, see [Dolstra et al. \(2004\)](#).

By leveraging these principles, **Nix** can effectively replace and integrate the functionalities of

tools like **renv** and **Docker** for R projects, or `requirements.txt` files and virtual environments for Python projects. Furthermore, **Nix** excels at building polyglot environments, which can seamlessly combine R, Python, Julia, a LaTeX distribution, or any of the numerous other tools available in **nixpkgs**. This enables the creation of a truly complete, project-specific, and deeply reproducible environment that can be used interactively for development or non-interactively for automated analysis. As long as the **nixpkgs** repository (or a suitable fork, as discussed in Section 4) remains accessible, the environment can be rebuilt reliably in the future.

Nix provides strong cross-platform support across major operating systems. It runs natively on Linux (x86_64) with full functionality and optimal integration. On macOS (Intel and Apple Silicon), however, reproducibility is less reliable due to dependencies on impure system frameworks and Xcode toolchains outside the Nix store. As a result, macOS users may need to update project pins after system or Xcode upgrades to restore builds. On Windows, **Nix** works well through the Windows Subsystem for Linux (WSL2), though graphical applications may need extra setup. **Nix** environments on WSL can be used interactively via VS Code or Positron for a smooth development experience. It is also important to note that **Nix** can be installed within a **Docker** image to deploy and reproduce environments reliably in containerized workflows.

While **Nix** is a powerful system, it comes with notable challenges. The steepest is its learning curve: **Nix** uses its own declarative language (also called Nix), which can be difficult for those new to functional or declarative paradigms. Complexity becomes most apparent when writing custom derivations or troubleshooting intricate build issues. Build times can also be long, especially for large environments, when pre-built binaries aren't available. Additionally, because **Nix** stores all package outputs immutably in the Nix store, disk usage tends to be higher than with traditional package managers. Finally, while **nixpkgs** is extensive, some new or niche packages may be missing or fail to build across platforms.

To make these capabilities more approachable for R users seeking reproducible, project-specific environments without learning the full Nix language, we created the **rix** and **rixpress** packages.

3. Reproducible development environments with Nix

As mentioned, **Nix** expressions are written in the Nix programming language, which is purely functional. Here is a simple example that creates a shell environment containing version 4.3.1 of R:

```
let
  pkgs = import (fetchTarball
    "https://github.com/NixOS/nixpkgs/archive/976fa336.tar.gz"
  ) {};
  system_packages = builtins.attrValues {
    inherit (pkgs) R;
  };
in
  pkgs.mkShell {
    buildInputs = [ system_packages ];
```

```
    shellHook = "R --vanilla";
}
```

In this expression, the `let` keyword is used to define variables. The variable `pkgs` imports the set of packages from the `nixpkgs` repository at the specified commit `976fa336`. The variable `system_packages` lists the packages to include in the environment; in this case, it is just the R programming language, along with all its dependencies and their transitive dependencies. The `mkShell` function then creates a development shell with the specified packages. The `shellHook` is set to `"R --vanilla"`, meaning that entering the shell automatically starts R in vanilla mode, ignoring any startup options.

This expression can be saved in a file called `default.nix`. The environment can then be built on a system with **Nix** installed using the `nix-build` command.² Once the build completes, the user can enter the interactive shell with `nix-shell`, which executes following steps:

- Nix reads the expression and identifies what needs to be installed;
- It checks if these packages already exist in `/nix/store`;
- If not, it downloads or builds them (with all dependencies);
- It creates a new shell environment and updates the `$PATH` to include the right binaries from `/nix/store`.
- The user is now in an interactive **Nix** shell session.

This shell contains all the packages specified in `default.nix` and can be used for development, similar to activating a virtual environment in the Python ecosystem.

Writing **Nix** expressions can be challenging for users unfamiliar with the Nix language. However, the ability to define a fully reproducible development environment in a single text file and then rebuild it anywhere is highly appealing. **rix** aims to lower the barrier to adoption of **Nix** for reproducibility.

rix provides the `rix()` function, which simplifies generating **Nix** expressions. It is available on CRAN and can be installed like any other R package. Additionally, it can bootstrap an R development environment on a system where R is not yet installed but **Nix** is available. This can be done by running (inside of a terminal):

```
$> nix-shell -I \
+   nixpkgs=https://github.com/rstats-on-nix/nixpkgs/tarball/2025-10-20 -p \
+   R rPackages.rix
```

(the `-I` flag allows one to pass a specific revision of `nixpkgs`, ensuring temporary shells are also reproducible).

This command opens a temporary R session with **rix** available.³ From there, users can generate new **Nix** expressions for building environments. For example, the following generates a `default.nix` file that installs R 4.3.1 along with the **dplyr** and **chronicler** packages:

```
R> library('rix')
```

²For installing **Nix**, we recommend the Determinate Systems installer: <https://determinate.systems/posts/determinate-nix-installer>

³`nix-shell -p` starts an interactive shell with the specified packages.

```
R> rix(r_ver = "4.3.1",
+     r_pkgs = c("dplyr", "chronicler"),
+     project_path = ".",
+     overwrite = TRUE)
```

rix can also handle more complex setups, and users can provide a date instead of a specific R version:

```
R> rix(date = "2025-10-20",
+     r_pkgs = c("rix", "dplyr", "chronicler", "AER@1.2-8"),
+     system_pkgs = c("quarto", "git"),
+     tex_pkgs = c(
+         "amsmath",
+         "framed",
+         "fvextra",
+         "environ",
+         "fontawesome5",
+         "orcidlink",
+         "pdfcol",
+         "tcolorbox",
+         "tikzfill"
+     ),
+     git_pkgs = list(
+         package_name = "fusen",
+         repo_url = "https://github.com/ThinkR-open/fusen",
+         commit = "60346860111be79fc2beb33c53e195f97504a667"
+     ),
+     ide = "positron",
+     project_path = ".",
+     overwrite = TRUE)
```

This call to `rix()` generates a `default.nix` file for a development shell that encapsulates a complete and reproducible research environment. It provides R and its packages (including **AER** at version 1.2-8), several TeXLive packages for L^AT_EX document authoring, development versions of **rix** and **fusen** pulled directly from GitHub at a specific commit, and the Positron editor. The reproducibility of this environment is guaranteed by pinning all components to a single point in time: the R packages are resolved from the CRAN snapshot of October 20, 2025, while all other system tools are fixed to the version of `nixpkgs` from that same date. Typing `nix-shell` in a terminal within the folder that contains this `default.nix` will drop the user into a new shell. It should be noted that it uses the `nixpkgs` from that same date. It is also possible to configure IDEs to dynamically load this new shell to provide the proper development tooling.

It is also possible to include Python and Julia packages using the `py_conf` and `j1_conf` arguments respectively. Julia packages work the same way as R packages: they are pinned to their versions as of the specified date.

Python packages, however, require special handling. Unlike CRAN or the Julia package registry, PyPI (the Python Package Index) contains over 500,000 packages with highly variable

quality—many lack proper build specifications, have broken dependencies, or are abandoned. Because of this heterogeneity, **Nix** maintainers cannot automatically package all of PyPI. Instead, they manually curate and package individual Python packages that meet quality standards. This means some Python packages, or specific versions, may not be available through **Nix**.

For projects requiring Python packages not yet available in **Nix**, **rix** supports integration with **uv**, a modern Python package manager. This allows users to combine Nix’s system-level reproducibility (Python interpreter, system libraries) with **uv**’s comprehensive package availability. While this hybrid approach sacrifices some of Nix’s deterministic build guarantees for Python packages, it provides a pragmatic solution when the needed packages are unavailable in **nixpkgs**.

Users can include **uv** as a system package:

```
R> rix(date = "2025-10-20",
+   r_pkgs = c("rix", "dplyr", "chronicler"),
+   system_pkgs = c("uv"),
+   project_path = ".",
+   overwrite = TRUE)
```

This hybrid approach offers complementary strengths: **Nix** manages the Python interpreter, system libraries (e.g., **GDAL**, **HDF5**), and R packages with full determinism, while **uv** handles Python packages through its own lockfile mechanism (**uv.lock**). Within the Nix environment, users run standard **uv** commands (e.g., **uv pip install pandas==2.0.0**) to install Python packages. The **uv** lockfile records exact package versions and their hashes, providing reproducibility for the Python package layer.

While this approach doesn’t achieve Nix’s level of build-time determinism for Python packages (**uv** still downloads wheels or builds packages at install time rather than using pre-built, content-addressed artefacts) it provides a pragmatic balance: comprehensive Python package availability with reasonable reproducibility guarantees. For polyglot projects, this means R and system-level dependencies remain fully reproducible through **Nix**, while Python packages gain the reproducibility that **uv** can provide.

rix can generate **Nix** expressions even if **Nix** is not installed on the system. This is useful for continuous integration and continuous deployment (CI/CD) workflows on platforms such as GitHub Actions. For instance, the repository containing the source code for this article⁴ uses GitHub Actions to compile the paper. Each time a push is made to the master branch, a runner installs **Nix**, generates the environment from the hosted **default.nix** file, and compiles the paper using **Quarto** within the reproducible environment. This ensure that *exactly* the same environment is used on the author’s computer and on the CI/CD without any additional, platform-specific, configuration.

Instead of first entering a **Nix** shell, it is also possible to run a program directly from the environment:

```
cd /path/to/project/ && nix-shell default.nix --run "Rscript analysis.R"
```

⁴https://github.com/b-rodrigues/rix_paper

This command runs `Rscript` and executes the `analysis.R` script, which in this example should be located in the same directory as `default.nix`.

4. The `rstats-on-nix` fork of `nixpkgs`

As explained earlier, **Nix** uses expressions from the `nixpkgs` GitHub repository to build software. However, when generating expressions with `rix`, the fork `rstats-on-nix/nixpkgs` is used instead.

Using a fork offers several advantages. First, it provides flexibility that the official `nixpkgs` repository cannot always accommodate. **Nix** is primarily the package manager for the NixOS Linux distribution, and governance and technical choices made upstream can limit what `rix` aims to provide.

For instance, while **Nix** can theoretically support multiple versions (or *variants*) of the same package, in practice maintainers cannot provide several variants for all R packages, given the size of the ecosystem (over 20,000 CRAN and Bioconductor packages). This makes it difficult to install a specific version of an R package not included in a particular `nixpkgs` commit. With `rix`, users can install a specific package version from source, e.g.:

```
R> rix(..., r_pkgs = "dplyr@1.0.7", ...)
```

However, installing from source might fail, especially if the package needs to be compiled.

Additionally, updating the full R package set on **Nix** daily is impractical. While CRAN and Bioconductor update daily, the R packages in `nixpkgs` are only updated with new R releases. This limitation is due to **Nix**'s governance as a Linux distribution package manager.

The `rstats-on-nix` fork allows me to circumvent these limitations. For example, it provides daily snapshots of CRAN. Each day, the R package set is updated and committed to a dated branch using GitHub Actions. Users can select a specific date with:

```
R> rix(date = "2024-12-14", ...)
```

We strive to provide an available date per week: each Monday, a GitHub Action tests popular packages on Linux and macOS, and only if all tests succeed is the date added to the list of available dates in `rix`. This ensures users can reliably install packages, and allows me to backport fixes if needed. For example, when RStudio was temporarily broken due to a dependency issue (`boost`), a pull request was submitted to the official `nixpkgs` repository. I backported the fix to the `rstats-on-nix` fork, making RStudio available to users of `rix` earlier than upstream, as merging PRs in the official repository can take some time.

We have backported fixes to the `rstats-on-nix nixpkgs` fork as far back as March 2019. The process involves checking out a `nixpkgs` commit on the selected date, updating the R package set using Posit CRAN and Bioconductor snapshots, backporting fixes, and ensuring popular packages work on both x86-linux (including WSL2) and aarch64-darwin (Apple Silicon). These changes are committed to a dated branch in `rstats-on-nix/nixpkgs`. Users can see all available dates with `rix::available_dates()`.

A drawback of forking `nixpkgs` is that backported packages are not included upstream and thus are not prebuilt by Nix's CI platform, Hydra. Users may need to build many packages

from source, which can be time-consuming. To mitigate this, we provide a binary cache sponsored by [Cachix](#), complementing the public Nix cache. Instructions for using Cachix are in **rix**'s documentation. Using the cache significantly speeds up installations, as prebuilt packages are downloaded rather than compiled.

The fork also allows me to catch issues (such as packages' builds breaking) early on, and prepare fixes that can then be contributed upstream.

While the reliance on the **rstats-on-nix** fork might initially raise sustainability concerns, this risk should be contextualized. The fork serves a focused purpose: it updates CRAN and Julia package registries daily and backports fixes when builds break. Importantly, users are never locked into this fork—they retain the ability to point directly to any commit in the upstream **nixpkgs** repository, maintaining full access to the official **Nix** ecosystem. The fork introduces no fundamental architectural changes or proprietary extensions; it functions as a convenience layer providing tested package sets and dated snapshots. Should maintenance cease, users could seamlessly transition to upstream **nixpkgs** commits, though with trade-offs: less frequent R and Julia package updates. The fork's value lies not in creating dependency, but in reducing friction through testing infrastructure and curated “known-good” dates. Thus, the fork represents an enhancement to user experience rather than a critical single point of failure—the underlying **Nix** infrastructure and official **nixpkgs** repository remain the foundation, ensuring long-term viability of environments created with **rix**.

5. Orchestrating the workflow with **rixpress**

Defining a reproducible environment with **rix** addresses the first major challenge of reproducibility. The second challenge is reliably and efficiently executing the analysis workflow within that environment, which is the role of the **rixpress** package (or **ryxpress** for Python users).

As mentioned in the introduction, a build automation tool like **targets** is invaluable for managing complex analyses. It tracks dependencies between code and data, caches results, and only recomputes steps that have changed. One can run a **targets** pipeline inside a Nix environment to make it reproducible. However, this approach has limitations: the entire pipeline must run in a single environment, and orchestrating steps across different languages (e.g., R and Python) requires manual handling via packages like **reticulate**.

rixpress overcomes these limitations by using **Nix** not just as a package manager, but as the build automation engine itself. In a **rixpress** pipeline, each step is defined as a **Nix** derivation, providing two key benefits:

1. True Polyglot Pipelines: Each step can have its own **Nix** environment. A Python step can run in a pure Python environment, an R step in an R environment, and a Quarto rendering step in yet another, all within the same pipeline.
2. Deep Reproducibility: Each step is a hermetically sealed **Nix** derivation whose output is cached in the **Nix** store based on the hash of all its inputs. All artefacts are direct children of the computational environment (because the computational environment is actually a dependency of the artefacts), ensuring they are rebuilt if the environment changes, keeping environment and outputs always in sync.

To illustrate these capabilities, we present a polyglot pipeline that simulates a Real Business Cycle (RBC) model in Julia, uses the resulting data to train an **XGBoost** forecasting model in Python, visualises the results in R with **ggplot2**, and finally compiles a **Quarto** report. The code for this example can be found in the [following repository](#)⁵, with additional details provided in the Appendix of this manuscript.

For new projects, **rixpress** provides `rxp_init()`, which generates the necessary boilerplate structure including a pipeline definition file and template helper scripts. Users then define their pipeline in the script `gen-pipeline.R` using functions inspired by **targets**, and the computational environment the pipeline will be executed in the `gen-env.R` script. If functions are needed for pipeline steps, these are to be specified, in language-specific helper scripts: `functions.jl` for Julia, `functions.py` for Python, and `functions.R` for R. This separation keeps the pipeline definition clean and readable while encapsulating the scientific logic in appropriate language-native files. Importantly, these helper scripts are standard language files—they contain no **rixpress**-specific code and can be reused in other projects or executed independently outside of **Nix**.

The following example demonstrates how **rixpress** orchestrates a granular, multi-step data science workflow that crosses language boundaries. (Complete setup instructions and a detailed walkthrough are provided in Appendix)

```
R> library('rixpress')

R> pipeline_steps <- list(
+   # STEP 0: Define RBC Model Parameters in Julia
+   rxp_jl(alpha, 0.3),
+   rxp_jl(beta, 1 / 1.01),
+   # ... (other parameters omitted for brevity) ...
+
+   # STEP 1: Julia - Simulate the RBC model
+   rxp_jl(
+     name = simulated_rbc_data,
+     expr = "simulate_rbc_model(alpha, beta, delta, rho, sigma, sigma_z)",
+     user_functions = "functions/functions.jl",
+     encoder = "arrow_write"
+   ),
+
+   # STEP 2.1: Python - Prepare features
+   rxp_py(
+     name = processed_data,
+     expr = "prepare_features(simulated_rbc_data)",
+     user_functions = "functions/functions.py",
+     decoder = "pyarrow.feather.read_feather"
+   ),
+
+   # STEP 2.2: Python - Split data (X_train, y_train, etc.)
+   rxp_py(name = X_train, expr = "get_X_train(processed_data)", ...),
```

⁵https://github.com/b-rodrigues/rixpress_demos/tree/master/case-study

```

+ # ... (other data splits omitted for brevity) ...
+
+ # STEP 2.3: Python - Train the XGBoost model
+ rxp_py(
+   name = trained_model,
+   expr = "train_model(X_train, y_train)",
+   user_functions = "functions/functions.py"
+ ),
+
+ # STEP 2.4: Python - Make predictions and format results
+ # ... (prediction and formatting steps omitted for brevity) ...
+ rxp_py(
+   name = final_predictions_df,
+   expr = "format_results(y_test, model_predictions)",
+   user_functions = "functions/functions.py",
+   encoder = "save_arrow"
+ ),
+
+ # STEP 3: R - Visualise the predictions
+ rxp_r(
+   name = output_plot,
+   expr = plot_predictions(final_predictions_df),
+   user_functions = "functions/functions.R",
+   decoder = arrow::read_feather
+ ),
+
+ # STEP 4: Quarto - Compile the final report
+ rxp_qmd(
+   name = final_report,
+   qmd_file = "readme.qmd"
+ )
+)

```

Generate the 'pipeline.nix' file from the R list

```
R> rxp_populate(pipeline_steps, build = TRUE)
```

Computation steps, referred to as *derivations*, are defined using the functions `rxp_r()`, `rxp_jl()`, and `rxp_py()`, corresponding to R, Julia, and Python environments, respectively. Variants of these functions suffixed with `_file()` designate *import derivations*, which introduce external data into the pipeline. It is important to note that any data imported in this way is stored within the **Nix** store, ensuring reproducibility but also persisting the data in the build environment. The function `rxp_qmd()` is used to compile a **Quarto** document (there is also `rxp_rmd()` which serves the same purpose for **R Markdown** documents).

The `rxp_populate()` function translates this R list into a `pipeline.nix` file, which declaratively defines the entire workflow. Data flows from derivation to derivation by being serialised into the efficient and language-agnostic Arrow format using the pair of `encoder/decoder` functions. Other universal formats, such as `csv` or `json` could have been used.

As a sidenote: Python users who wish to use **ryxpress** define pipelines using the same R-based DSL shown above. This design choice keeps the pipeline definition language consistent across both R and Python ecosystems. **ryxpress** calls R and **rixpress** under the hood to generate and build the `pipeline.nix` file. To use **ryxpress**, one should simply add this package to the `py_conf` argument of `rix()`.

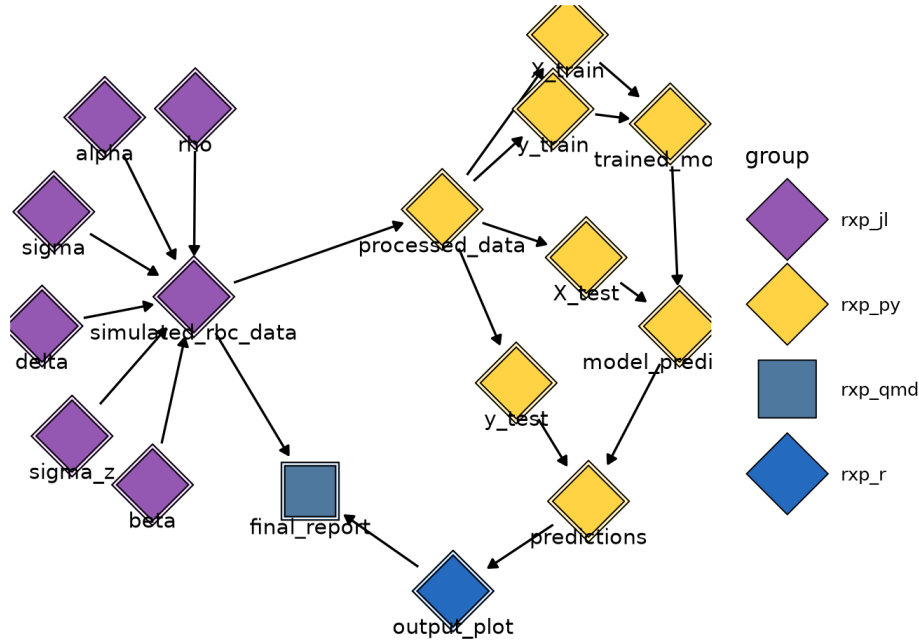


Figure 2: Graphical representation of the polyglot pipeline. Purple nodes are **Julia**, yellow nodes are **Python**, light blue nodes are **R**, and dark blue nodes are **Quarto** derivations.

The package can then generate a visual representation of the pipeline’s directed acyclic graph by running `rxp_dag()`, as can be seen in Figure 2.

To execute the pipeline, one can either set `build = TRUE` in `rxp_populate()` or call `rxp_make()` separately. **Nix** executes each step in order, building dependencies as needed. Outputs are cached, so subsequent runs only recompute steps with changed inputs or code. This provides the efficiency of **targets** with polyglot support and bit-for-bit reproducibility. Artefacts can be inspected interactively in R using `rxp_read("artefact_name")` or `rxp_load("artefact_name")`.

rixpress also includes several additional features not covered here for brevity. As mentioned previously, it is also possible to configure popular IDEs to work interactively and seamlessly with both **rix** and **rixpress**, enabling a smooth, reproducible workflow from within the development environment. Detailed setup instructions are provided in the vignettes of both packages.

6. A Comparison of Pipeline Paradigms: Imperative vs. Declarative

To illustrate the practical benefits of our framework, we implemented the RBC pipeline exam-

ple from Section 5 using two distinct paradigms. The first is a traditional, imperative stack combining Docker for environment isolation, Make for workflow orchestration, and language-specific package managers. The second is our proposed declarative framework using `rix` and `rixpress`. While both approaches successfully execute the same polyglot analysis to produce identical outputs, their comparison illuminates fundamental differences in conceptual models, system complexity, and the nature of the reproducibility guarantees they provide.

6.1. The Imperative Composition of Specialized Tools

The traditional approach represents a composition of specialized tools, each solving part of the reproducibility puzzle. This paradigm is entirely imperative; the researcher must explicitly script the procedural steps for both environment construction and workflow execution.

The codification of the environment is a procedural build process specified in a `Dockerfile`. Successfully authoring such a file necessitates that the researcher act as a system administrator, manually scripting a sequence of shell commands to install each component. This includes managing `apt` repositories for R, invoking external installers like `uv` and `rig` for Python and R runtimes, and using `curl` and `tar` for Julia. Subsequently, separate commands must be run to install packages within each language's ecosystem. The procedural verbosity and intertwined concerns of this approach are reflected in the resulting `Dockerfile`, which spans 108 lines to define the complete environment.

Here is a sketch of this `Dockerfile`:

```
# Add R repository and install specific version
RUN apt-get update && apt-get install -y software-properties-common
RUN add-apt-repository ppa:...
RUN curl -L https://rig.r-pkg.org/... | sh
RUN rig add 4.5.1

# Install Python with uv
RUN curl -LsSf https://astral.sh/uv/install.sh | sh
RUN uv python install 3.13

# Download and extract Julia
RUN curl -fsSL https://julialang-s3.julialang.org/... -o julia.tar.gz
RUN tar -xzf julia.tar.gz -C /opt/

# Install packages for each language separately
RUN echo 'options(repos = c(CRAN = ...))' > /root/.Rprofile
RUN Rscript -e 'install.packages(...)'

# Install Python packages using uv with specific versions for reproducibility.
RUN echo "pandas==2.3.3" > /tmp/requirements.txt && \
    echo "scikit-learn==1.7.2" >> /tmp/requirements.txt && \
    # ... more packages ...

RUN uv pip install --no-cache -r /tmp/requirements.txt && \
```

```
rm /tmp/requirements.txt

# Install specific versions of Julia packages for reproducibility
RUN julia -e 'using Pkg; \
  Pkg.add(name="Arrow", version="2.8.0"); \
  # ... more packages ...'
```

This procedural philosophy extends from environment definition to workflow orchestration. A **Makefile** defines the pipeline as a directed acyclic graph of file dependencies. While robust, this model requires the researcher to manually encode the dependency graph, tightly coupling the analytical logic to specific file paths. Furthermore, since **make** is agnostic to the content of these files, it provides no mechanism to enforce type or schema consistency between steps.

A direct consequence of this file-based orchestration model is the proliferation of wrapper scripts. Each step in the **Makefile** invokes a script containing a significant amount of procedural boilerplate unrelated to the scientific logic. This includes parsing command-line arguments, handling data serialization and deserialization, and managing file paths. The core analytical functions are thus obscured by this necessary but distracting scaffolding.

The cumulative effect is a system of high structural complexity. The traditional implementation requires nine separate files to fully specify the environment, workflow, and analysis. This includes the **Dockerfile**, **Makefile**, three wrapper scripts for orchestration, and three **functions** files containing the core logic, in addition to the final report. The cognitive overhead is substantial, as a user must comprehend the conventions and interactions of this entire toolchain.

6.2. A Unified Declarative Framework

The Nix-based framework inverts this imperative model. Instead of scripting the procedural construction of the environment and workflow, the researcher declares the desired end-state.

Environment specification is reduced to a single declarative statement in the **gen-env.R** script. The researcher specifies *what* the environment should contain, leaving the procedural implementation of *how* to build it entirely to the Nix toolchain. The **gen-env.R** script replaces the 108-line **Dockerfile**, abstracting away all system administration details. **rix** translates this high-level specification into a formal Nix expression that deterministically resolves all transitive dependencies for the complete polyglot environment.

```
library(rix)
rix(
  date = "2025-10-14",
  r_pkgs = c("ggplot2", "dplyr", "arrow"),
  jl_conf = list(jl_version = "lts", ...),
  py_conf = list(py_version = "3.13", ...),
  ...
)
```

Workflow orchestration undergoes a similar conceptual shift from procedural to declarative. The **gen-pipeline.R** script defines the pipeline not as a set of file rules, but as a graph of data

objects. Dependencies are not manually specified but are automatically inferred by **rixpress** from the functional relationships in the code. This object-centric model renders the procedural wrapper scripts unnecessary. The scientific logic is expressed as a set of pure functions that accept data structures as inputs and return them as outputs, free of any boilerplate for file I/O or command-line parsing.

```
# functions/functions.py
def train_and_predict_output(simulated_df: pd.DataFrame) -> pd.DataFrame:
    # ... scientific logic ...
    return format_results(predictions)
```

This declarative approach significantly reduces system complexity. The entire implementation requires only six files: two for the high-level environment and pipeline declarations, three containing the pure scientific functions, and the final report. The reduction is not merely quantitative; the remaining files are conceptually simpler, containing only analytical logic and specifications, with all procedural scaffolding abstracted away by the framework.

6.3. Reproducibility Guarantees: Space and Time

The critical distinction between the two paradigms lies in the nature of their reproducibility guarantees. The traditional stack, anchored by Docker, provides strong **spatial reproducibility**—the ability to execute the analysis identically across different machines. However, it is vulnerable to **temporal drift**, a concept that distinguishes reproducibility across space from consistency over time (?). This temporal vulnerability arises from several sources, including the resolution of mutable base image tags (e.g., **ubuntu:24.04**) and the potential for language-level package managers to resolve dependencies differently over time.

The Nix-based approach is architected to solve this problem of temporal drift. By pinning the entire software ecosystem to a single, immutable revision of the **nixpkgs** repository, the framework provides strong guarantees of **temporal reproducibility**. This claim is supported by large-scale empirical evidence; a study by ?, which rebuilt over 700,000 historical packages, confirmed that Nix’s functional model achieves a bit-for-bit reproducibility rate exceeding 90% and a rebuildability rate over 99% across a multi-year timeframe.

6.4. Summary of Differences

Dimension	Traditional Stack	Nix-Based Framework
Paradigm	Imperative composition of specialized tools	Declarative specification in unified framework
Environment setup	Step-by-step shell commands	High-level specification
Workflow definition	File-based rules in Make syntax	Object-based graph in R syntax
Scientific code	Wrapped in CLI parsing and file I/O boilerplate	Pure functions with no plumbing
File count	9 files (4 for orchestration and plumbing)	6 files (2 for orchestration)

Dimension	Traditional Stack	Nix-Based Framework
Expertise required	Linux system administration, Make syntax	R programming
Spatial reproducibility	Strong (via Docker isolation)	Strong (via Nix isolation)
Temporal reproducibility	Moderate (vulnerable to drift)	Strong (validated empirically (Malka <i>et al.</i> 2025))

This declarative rigor, however, is not without pragmatic costs. Our validation on GitHub Actions shows that the imperative Docker-based runner completes the entire pipeline in approximately three minutes, whereas the declarative Nix-based runner takes around five minutes. This difference reflects the trade-off between the upfront computational cost of Nix’s rigorous dependency resolution and the faster, but more brittle, imperative builds of the traditional approach.

7. Conclusion

Many tools exist to improve reproducibility, but **Nix** stands out because it deploys complete software environments *closed under the “depends on” relation*: it installs not only a package, but all its dependencies and their dependencies. This makes **Nix** particularly powerful for reproducible research.

However, solving such a complex problem makes **Nix** a complex tool. With **rix**, we aim to make **Nix** more accessible to R users by providing a familiar interface and workflow. By building reproducible development shells with **Nix**, researchers can accommodate a wide range of use cases: running scripts and pipelines, developing interactive **shiny** applications, or serving **plumber** APIs.

Furthermore, **rixpress** extends this reproducibility to entire analysis pipelines. By leveraging **Nix** as a build automation engine, **rixpress** allows polyglot workflows where each step runs in its own hermetically sealed environment. This ensures deep reproducibility, efficient caching, and seamless orchestration of polyglot analyses. Most importantly, it provides native polyglot support without the manual coordination required by containerization approaches, making complex multi-language workflows accessible to researchers without systems administration expertise.

Adopting **Nix**-based workflows represents a significant paradigm shift, and researchers should approach it incrementally. A reasonable pathway begins with using **rix** to generate reproducible environments for new projects, while maintaining existing workflows for ongoing work. This allows learning **Nix** concepts without risking active research. As comfort grows, researchers can experiment with **rixpress** for smaller analyses before tackling complex polyglot pipelines. Institutions can support adoption by providing workshops, maintaining institutional binary caches with private packages, and designating local experts. The investment pays increasing returns: early projects require substantial learning, but subsequent projects leverage accumulated knowledge and reusable environment definitions. For research groups, having one member develop **Nix** expertise can enable the entire team to benefit from repro-

ducible environments without everyone climbing the learning curve. Ultimately, the decision to adopt these tools depends on project timescales, reproducibility requirements, and institutional context—but the framework exists for those who need it.

While **rix** and **rixpress** (and **ryxpress**) provide significant steps forward, it is important to acknowledge current limitations. Debugging complex **rixpress** pipelines remains more difficult than debugging **targets** workflows, particularly when failures occur deep in the **Nix** build process. The visualization capabilities, while useful, lack features like automatic detection of outdated derivations that make tools like **targets** particularly developer-friendly. Cross-platform reproducibility, while generally strong, has known limitations on macOS due to system framework dependencies. These are not insurmountable problems—they represent areas for continued development and community contribution. Researchers adopting these tools should do so with realistic expectations: they provide powerful capabilities for deep reproducibility and polyglot workflows, but require patience during the learning phase and may occasionally present challenges that simpler tools avoid.

Acknowledgments

We thank the rOpenSci reviewers and contributors who provided valuable feedback on the development of **rix** and **rixpress**. In particular, we are grateful to David Watkins and Jacob Wujciak-Jens for their reviews of **rix**, and to William Landau and Anthony Martinez for their reviews of **rixpress**. We also acknowledge the contributions of Richard J. Acton, Jordi Rosell, Elio Campitelli, László Kupcsik, and Michael Heming for **rix**. Their expertise and feedback greatly improved the quality and usability of these packages. We would also like to thank Pol Dellaiera and Edvin Syk for providing feedback on the manuscript.

References

- Bezanson J, Edelman A, Karpinski S, Shah VB (2017). “Julia: A fresh approach to numerical computing.” *SIAM review*, **59**(1), 65–98. URL <https://doi.org/10.1137/141000671>.
- Bhandari Neupane J, Neupane RP, Luo Y, Yoshida WY, Sun R, Williams PG (2019). “Characterization of Leptazolines A–D, Polar Oxazolines from the Cyanobacterium *Leptolyngbya* sp., Reveals a Glitch with the “Willoughby–Hoye” Scripts for Calculating NMR Chemical Shifts.” *Organic Letters*, **21**(20), 8449–8453. doi:10.1021/acs.orglett.9b03332.
- Boettiger C, Eddelbuettel D (2017). “An Introduction to Rocker: Docker Containers for R.” *The R Journal*, **9**(2), 527–536. doi:10.32614/RJ-2017-065.
- De Paoli B (2009). “Slides 1: The RBC Model, Analytical and Numerical solutions.” https://personal.lse.ac.uk/depaoli/RBC_slides1.pdf. Lecture Slides.
- Dellaiera P (2024). “Reproducibility in Software Engineering.” doi:10.5281/zenodo.12666898.
- Dolstra E, De Jonge M, Visser E (2004). “**Nix**: A Safe and Policy-Free System for Software Deployment.” In *18th Large Installation System Administration Conference*, pp. 79–92.

- Fay C, Guyader V, Parry J, Rochette S (2024). **dockerfiler**: *Easy Dockerfile Creation from R*. R package version 0.2.5, URL <https://thinkr-open.github.io/dockerfiler/>.
- hong Chan C, Schoch D (2023). “**rang**: Reconstructing Reproducible R Computational Environments.” *PLOS ONE*, **18**(6), e0286761. doi:10.1371/journal.pone.0286761.
- Landau WM (2021). “The **targets** R Package: A Dynamic Make-like Function-Oriented Pipeline Toolkit for Reproducibility and High-Performance Computing.” *Journal of Open Source Software*, **6**(57), 2959. doi:10.21105/joss.02959.
- Malka J, Zacchiroli S, Zimmermann T (2024). “Reproducibility of Build Environments through Space and Time.” In *Proceedings of the 2024 ACM/IEEE 44th International Conference on Software Engineering: New Ideas and Emerging Results*, ICSE-NIER’24, p. 97–101. ACM. doi:10.1145/3639476.3639767. URL <http://dx.doi.org/10.1145/3639476.3639767>.
- Malka J, Zacchiroli S, Zimmermann T (2025). “Does Functional Package Management Enable Reproducible Builds at Scale? Yes.” In *22nd International Conference on Mining Software Repositories*. Ottawa, Canada. URL <https://hal.science/hal-04913007>.
- McDermott GR (2021). “Skeptic Priors and Climate Consensus.” *Climatic Change*, **166**(1-2), 7. doi:10.1007/s10584-021-03114-6.
- Peng RD (2011). “Reproducible Research in Computational Science.” *Science*, **334**(6060), 1226–1227. doi:10.1126/science.1213847.
- R Core Team (2021). *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria. URL <https://www.R-project.org/>.
- R Infrastructure (2023). **rig**: *The R Installation Manager*. URL <https://github.com/r-lib/rig>.
- Simonsohn U, Gruson H (2023). *groundhog: Version-Control for CRAN, GitHub, and GitLab Packages*. R package version 3.1.2, <https://github.com/CredibilityLab/groundhog>, URL <https://groundhogr.com/>.
- Van Rossum G, Drake FL (2009). *Python 3 Reference Manual*. CreateSpace, Scotts Valley, CA. ISBN 1441412697.

8. Appendix

8.1. Reproducing this paper

The source code of this paper is hosted on GitHub and can be found at following [link](#)⁶. The paper can be easily compiled by running the following command:

```
$> nix-shell --run "quarto render paper.qmd --to jss-pdf"
```

⁶https://github.com/b-rodrigues/rix_paper

The `default.nix` file defines the exact computational environment required to compile the manuscript, ensuring full reproducibility of the build process. To guarantee the reproducibility of the manuscript, it is automatically recompiled via GitHub Actions upon each commit, using the same **Nix** environment. An HTML version is additionally deployed to GitHub Pages, providing an accessible format for viewing on smaller devices at this [link](#)⁷. The PDF version of the manuscript can be found at [link](#)⁸.

8.2. A Complete Polyglot Example with rixpress

This appendix provides a conceptual walk-through of the polyglot pipeline example discussed in Section 5. The pipeline simulates a Real Business Cycle (RBC) model in Julia, trains an **XGBoost** model in Python, visualises the results in R, and compiles a final report with **Quarto**.

The complete, runnable source code for this example is available in the paper’s GitHub repository. A shell script, `run_polyglot_example.sh`, is provided to automate the entire process from environment creation to final output, as described at the end of this section⁹.

Project Structure and Components

The project is organized into a set of scripts, each with a distinct responsibility. The core logic is separated from the environment definition and pipeline orchestration.

- **functions/**: This directory contains the helper scripts with the core analytical code for each language.
 - `functions.jl`: The Julia script for the economic simulation.
 - `functions.py`: The Python script for the machine learning workflow.
 - `functions.R`: The R script for data visualisation.
- `gen-env.R`: An R script that defines the reproducible computational environment.
- `gen-pipeline.R`: The main R script that defines and orchestrates the entire polyglot pipeline.

Step 1: The Environment Definition

The foundation of the project is the reproducible environment, defined in `gen-env.R`. This script uses the `rix()` function to programmatically generate a `default.nix` file. This file serves as a complete blueprint for the computational environment, specifying:

- The exact versions of R, Python, and Julia.
- A list of required packages for each language (e.g., **ggplot2** for R, **xgboost** for Python, **DataFrames** for Julia).
- System-level dependencies like **Quarto**.

Crucially, the entire environment is pinned to a specific date (2025-10-14), ensuring that anyone who builds the environment, now or in the future, will get the exact same software versions, guaranteeing reproducibility.

⁷https://b-rodrigues.github.io/rix_paper/

⁸https://b-rodrigues.github.io/rix_paper/paper.pdf

⁹https://github.com/b-rodrigues/rix_paper/blob/master/repro-script/run_polyglot_example.sh

By dropping into a temporary shell using the following command:

```
$> nix-shell -I \
+ nixpkgs=https://github.com/rstats-on-nix/nixpkgs/tarball/2025-10-20 -p \
+ R rPackages.rix
```

it is possible to generate the `default.nix` by sourcing `gen-env.R`. Then, one leaves this temporary shell, and builds the environment using the command `nix-shell`, which also drop the user into the development shell.

Step 2: The Core Analytical Logic

The scientific logic for each stage of the pipeline is encapsulated in the separate helper scripts within the `functions/` directory.

- The RBC Model Simulation (`functions.jl`): This Julia script contains a single pure function that implements the state-space solution to the RBC model, based on [De Paoli \(2009\)](#). It takes the model's economic parameters as inputs and returns the simulated time-series data as a **DataFrame**.
- The **XGBoost** Forecasting Model (`functions.py`): This Python script handles the complete machine learning workflow through a series of modular functions. Its responsibilities include feature engineering (creating lagged variables), splitting the data into training and testing sets, training the **XGBoost** model, and generating predictions.
- The Visualisation (`functions.R`): This R script contains a function that uses **ggplot2** to create the final visualisation. It is designed to take the data frame of actual and predicted values produced by the Python script and generate a plot comparing the two series.

Step 3: Orchestrating the Pipeline

The entire workflow is defined and orchestrated by `gen-pipeline.R`. This script acts as the master plan, using functions from the **rixpress** package to define each computational step as a *derivation*. It declaratively outlines the dependencies between steps:

1. It begins by defining the RBC model parameters in Julia.
2. It specifies that the RBC simulation in `functions.jl` depends on these parameters.
3. It then defines the series of Python steps (feature preparation, training, prediction), making each one dependent on the output of the previous one. The first Python step is explicitly made dependent on the output from the Julia simulation. **rixpress** handles the passing of data between languages, in this case using the **Arrow** file format for efficiency.
4. Finally, it defines the **R visualisation** step, which depends on the final predictions from the Python model, and a **Quarto rendering** step that depends on the generated plot.

When this script is run in the development shell (by executing `source("gen-pipeline.R")`), **rixpress** translates the declared pipeline into a master Nix expression that **Nix** can execute, automatically handling the caching of results and re-running only the necessary steps if a piece of code or data changes.

To build the pipeline from an interactive Python session, one would execute the following lines:

```
Python> from ryxpress import rxp_make  
Python> rxp_make()
```

Running the Project

This entire example can be executed by running the `run_polyglot_example.sh` script available in the root of the paper's repository. This script automates the full process: 1. It first executes `gen-env.R` inside a temporary **Nix** shell to build the `default.nix` file. 2. It then executes `gen-pipeline.R` inside the newly defined environment. This triggers **Nix** to run the entire polyglot pipeline in the correct order.

Upon completion, the script will have generated all intermediate artefacts and the final `readme.html` report found in the `pipeline-output` folder containing the visualisation.

Affiliation:

Bruno Rodrigues
Department of Statistics
18, Montée de la Pétrusse
Luxembourg Luxembourg
E-mail: bruno@brodrigues.co
URL: <https://www.brodrigues.co>

Philipp Baumann
Data and Analytics
7, Länggassstrasse
Bern Switzerland
E-mail: baumann-philipp@protonmail.com